

5 RAG Mistakes Indian Engineers Ship to Production.

*Costing them ₹50K – ₹2L per month.
Easily.*

SAVE THIS ↓

Default chunk size = 500 tokens for everything.

WRONG FOR

Legal docs, code, research papers

RIGHT APPROACH

Chunk by semantic boundary, not token count.

TOOL

LangChain RecursiveCharacterTextSplitter
with custom separators per content type

Using multilingual-e5 for Hindi retrieval.

-18% *recall on Hindi queries
vs. IndicBERT*

FIX

**AI4Bharat IndicBERT or Bhashini
Indic embeddings.**

- Free
- Hosted on HuggingFace
- Drops into LangChain in 4 lines of code

No reranking.

Top-5 retrieved $\neq$ *Top-5 relevant.*

+15–25%

recall@5 jump

FIX

Add a reranker.

- BGE-reranker-v2-m3 — free on HuggingFace
- Cohere Rerank — paid, lowest latency
- 50ms latency cost

Worth it.

No retrieval evals.

“It works on the demo.”

is not a metric.

BUILD THIS

Eval set with 50 real user queries.

Track:

- **recall@5**
- **MRR**
- **hit rate**

Run weekly. The failures will surprise you.

Stuffing 50K tokens into context.

Context rot is real.

Performance drops past 32K tokens — even on 200K-context models.

MANTRA

**Retrieve aggressively.
Prune ruthlessly.**

*5–15 high-signal chunks
beats 50 low-signal ones.*

Every. Single. Time.

REAL EXAMPLE · BENGALURU FINTECH

Monthly RAG bill before:

₹2.4L

After fixing 2 of these:

₹38K

Same response quality. Same model.

We cover all 5
(plus 12 more)
in Modules 5–6.

GET THE FULL ROADMAP

17-page GenAI Engineering Roadmap

Free · No email gate

Link in comments ↓

WhatsApp Community + Discord links there too.

Sart Kumar · Co-founder, Netsetos

Building production GenAI training for Indian engineers.